

UC San Diego Extension Cloud Services for Machine Learning

Summer 2020

Homework#2

Date Given: July 6, 2020

Due Date: July 12, 2020

=====

Analyze the data source in 'kc-house-data.csv' file. This data source is a part of databases available in the public domain. This file contains 21,613 observations of real-estate properties of King county in Washington state. The data for the following 21 variables are provided.

1. id
2. date
3. price
4. bedrooms
5. bathrooms
6. sqft_living
7. sqft_lot
8. floors
9. waterfront
10. view
11. condition
12. grade
13. sqft_above
14. sqft_basement
15. yr_built
16. yr_renovated
17. zipcode
18. latitude
19. longitude
20. sqft_living15
21. sqft_lot15

Read the raw data source file 'kc-house-data.csv'. Build a linear regression model (as described in Problem#1 and Problem#2 on the next page) using the following variables.

Response Variable:

- price (numerical)

Predictor Variables:

- sqft_living (numerical)
- bedrooms (numerical)
- waterfront (categorical):
 - Levels of waterfront: 0 = no waterfront, 1 = waterfront
- condition(categorical)
 - Levels of condition: 1,2,3,4,5

Problem#1

Build a regression model using the following characteristics.

- Programming Language: Python
- Cloud Platform: Colab
- Package: Scikit-Learn

Verify that your regression equation is as follows.

$$\text{price} = 66,581.53 + 305.72 * \text{sqftliving} - 52,704.36 * \text{bedrooms} + 783,090.68 * \text{waterfront} \\ - 25,698.33 * \text{condition2} - 8,811.75 * \text{condition3} + 28,198.78 * \text{condition4} + 100,565.72 * \text{condition5}$$

Predict the price of a home with following characteristics:

- sqft_living = 3,000
- bedrooms = 4
- waterfront = No (0)
- condition = 4

Verify that the *Predicted price* = \$801,112.20

=====

Problem#2

Build a regression model using the following characteristics.

- Programming Language: None
- Cloud Platform: AutoML GCP

The procedure to build a regression model on GCP is as follows.

1. GCP/Storage
 - a. Create a Bucket in GCP
 - b. Region: us-central1(Iowa)
 - c. Upload Data file in that bucket
2. GCP/Table/Dataset
 - a. Import data in a GCP Dataset from the bucket: Takes time
3. GCP/Table/Model
 - a. Train the Model
 - i. Select Target Variable + Budget: Takes time
 - b. Evaluate the Model
 - c. Test & Use: Deploy the Model: Takes time
 - i. Prediction

Predict the price of a home with following characteristics:

- sqft_living = 3,000
- bedrooms = 4
- waterfront = No (0)
- condition = 4

The predicted value of the 'price' variable using GCP should be approximately equal to \$801,112.20. Also compute the 95% prediction interval of the response variable 'price'.

Building a regression model on GCP will cost a certain amount. Please check the GCP charges on your account before and after you complete this assignment. Make sure you do not deplete the \$300 credit you have on your account.

How to handle 'condition' categorical variable in Scikit-Learn:

The 'condition' variable is categorical with 5 levels. The values of this variable are 1,2,3,4,5. This does NOT mean that $5 > 4 > 3 > 2 > 1$. Since there are 5 levels of this variable, we need to replace the 'condition' variable with 4 (k-1) dummy (indicator) variables when performing the regression in Scikit-Learn.

When regression is performed in GCP/AutoML, we need to declare the 'condition' variable as categorical during the 'train' stage of model building. The GCP/AutoML does not need one-hot-encoded data as input. GCP/AutoML will perform the one-hot-encoding internally.

We must convert the 'condition' variable into 4 separate dummy variables using one-hot-encoding. The logic used for prediction is shown in the table below. The 'condition=1' will be our base condition. All values will be computed relative to 'condition=1'.

	Var: condition2	Var: condition3	Var: condition4	Var: condition5
Condition1 (Base)	0	0	0	0
Condition2	1	0	0	0
Condition3	0	1	0	0
Condition4	0	0	1	0
Condition5	0	0	0	1

Regression equation is as follows.

$$\text{price} = 66,581.53 + 305.72 * \text{sqftliving} - 52,704.36 * \text{bedrooms} + 783,090.68 * \text{waterfront} \\ - 25,698.33 * \text{condition2} - 8,811.75 * \text{condition3} + 28,198.78 * \text{condition4} + 100,565.72 * \text{condition5}$$

- This means that the price of a house with '**condition=2**' is \$25,698.33 **less** compared with the house with condition=0.
- This means that the price of a house with '**condition=5**' is \$100,565.72 **more** compared with the house with condition=0.

Now let us predict the price of the house using different value of the 'condition' categorical variable.

=====

Predict the price of a home with following characteristics:

- sqft_living = 3,000
- bedrooms = 4
- waterfront = No (0)
- **condition = 1**

$$\begin{aligned} \text{price} &= 66,581.53 + 305.72 * \text{sqftliving}(3,000) - 52,704.36 * \text{bedrooms}(4) + 783,090.68 \\ &\quad * \text{waterfront}(0) \\ \text{price} &= 772,913.4 \end{aligned}$$

=====

Predict the price of a home with following characteristics:

- sqft_living = 3,000
- bedrooms = 4
- waterfront = No (0)
- **condition = 2**

$$\begin{aligned} \text{price} &= 66,581.53 + 305.72 * \text{sqftliving}(3,000) - 52,704.36 * \text{bedrooms}(4) + 783,090.68 \\ &\quad * \text{waterfront}(0) - 25,698.33 * \text{condition2}(1) \\ \text{price} &= 747,215.1 \end{aligned}$$

=====

Predict the price of a home with following characteristics:

- sqft_living = 3,000
- bedrooms = 4
- waterfront = No (0)
- **condition = 3**

$$\begin{aligned} \text{price} &= 66,581.53 + 305.72 * \text{sqftliving}(3,000) - 52,704.36 * \text{bedrooms}(4) + 783,090.68 \\ &\quad * \text{waterfront}(0) - 8,811.75 * \text{condition3}(1) \\ \text{price} &= 764,101.7 \end{aligned}$$

=====

Predict the price of a home with following characteristics:

- sqft_living = 3,000
- bedrooms = 4
- waterfront = No (0)
- **condition = 5**

- $\text{price} = 66,581.53 + 305.72 * \text{sqftliving}(3,000) - 52,704.36 * \text{bedrooms}(4) + 783,090.68 * \text{waterfront}(0) + 100,565.72 * \text{condition5}(1)$
- $\text{price} = 873,479.2$

=====

Problem#1

Python Code



Housing_Data_Colab.ipynb ☆

File Edit View Insert Runtime Tools Help [All changes saved](#)

+ Code + Text

```
[1] import pandas as pd
import numpy as np
from sklearn import linear_model
```

```
[2] from google.colab import files
uploaded = files.upload()
```

Choose Files 00 kc_house_data.csv

- 00 kc_house_data.csv(application/vnd.ms-excel) - 2515206 bytes, last modified: 5/19/2016 - 100% done
Saving 00 kc_house_data.csv to 00 kc_house_data.csv

```
[ ] import io
#data = pd.read_csv(io.StringIO(uploaded['00 kc_house_data.csv'].decode('utf-8')))
data = pd.read_csv(io.BytesIO(uploaded['00 kc_house_data.csv']))
```

```
[3] data = pd.read_csv("00 kc_house_data.csv")
```

```
print(data.head())
print(data.tail())
data.shape
```

```

id      date      price  ...    long  sqft_living15  sqft_lot15
0  7129300520  20141013T000000  221900.0  ... -122.257      1340      5650
1  6414100192  20141209T000000  538000.0  ... -122.319      1690      7639
2  5631500400  20150225T000000  180000.0  ... -122.233      2720      8062
3  2487200875  20141209T000000  604000.0  ... -122.393      1360      5000
4  1954400510  20150218T000000  510000.0  ... -122.045      1800      7503
```

```
[5 rows x 21 columns]
```

```

id      date      price  ...    long  sqft_living15  sqft_lot15
21608  263000018  20140521T000000  ...    1530      1509
```

```
[5] predictedData = data[['sqft_living','bedrooms','waterfront','condition']]
```

```
[6] dummyWaterfront = pd.get_dummies(predictedData.waterfront,prefix='wf')
    print(dummyWaterfront.head())

    dummyCondition = pd.get_dummies(predictedData.condition,prefix='condition')
    print(dummyCondition.head())
```

```

    wf_0  wf_1
0      1    0
1      1    0
2      1    0
3      1    0
4      1    0
    condition_1  condition_2  condition_3  condition_4  condition_5
0              0            0            1            0            0
1              0            0            1            0            0
2              0            0            1            0            0
3              0            0            0            0            1
4              0            0            1            0            0

```

```
[7] merge=pd.concat([predictedData,dummyWaterfront,dummyCondition],axis='columns')
    merge.head()
```

```

    sqft_living  bedrooms  waterfront  condition  wf_0  wf_1  condition_1  condition_2  condition_3  condition_4  condition_5
0          1180         3          0          3     1     0             0             0             1             0             0
1          2570         3          0          3     1     0             0             0             1             0             0
2           770         2          0          3     1     0             0             0             1             0             0
3          1960         4          0          5     1     0             0             0             0             0             1
4          1680         3          0          3     1     0             0             0             1             0             0

```

```

    finalData = merge.drop(['waterfront','condition','wf_0','condition_1'],axis='columns')
    finalData.head()

```

```

    sqft_living  bedrooms  wf_1  condition_2  condition_3  condition_4  condition_5
0          1180         3     0             0             1             0             0
1          2570         3     0             0             1             0             0
2           770         2     0             0             1             0             0
3          1960         4     0             0             0             0             1
4          1680         3     0             0             1             0             0

```

```
[9] X = finalData
    print (type(X))
    print (X.shape)

    y = data['price']
    print (type(y))
    print (y.shape)
```

```
<class 'pandas.core.frame.DataFrame'>
(21613, 7)
<class 'pandas.core.series.Series'>
(21613,)
```

```
[10] linreg = linear_model.LinearRegression()
      linreg.fit(X, y)
      print (linreg.intercept_)
      print (linreg.coef_)
```

```
66581.52860950812
[ 3.05716451e+02 -5.27043605e+04  7.83090677e+05 -2.56983253e+04
 -8.81175295e+03  2.81987756e+04  1.00565722e+05]
```

```
linreg.predict([[3000,4,0,0,0,1,0]])
```

```
array([801112.21594909])
```

$$\text{price} = 66,581.53 + 305.72 * \text{sqftliving} - 52,704.36 * \text{bedrooms} + 783,090.68 * \text{waterfront} \\ - 25,698.33 * \text{condition2} - 8,811.75 * \text{condition3} + 28,198.78 * \text{condition4} + 100,565.72 * \text{condition5}$$

Predict the price of a home with following characteristics:

- sqft_living = 3,000
- bedrooms = 4
- waterfront = No (0)
- condition = 4

Predicted price = 801,112.2

Problem#2

Build a regression model using the following characteristics.

- Programming Language: None
 - Cloud Platform: AutoML GCP
4. GCP/Storage
 - a. Create a Bucket in GCP
 - b. Region: us-central1(Iowa)
 - c. Upload Data file in that bucket
 5. GCP/Table/Dataset
 - a. Import data in a GCP Dataset from the bucket: Takes time
 6. GCP/Table/Model
 - a. Train the Model
 - i. Select Target Variable + Budget: Takes time
 - b. Evaluate the Model
 - c. Test & Use: Deploy the Model: Takes time
 - i. Prediction

Predict label price	Prediction result 771,622.313	95% prediction interval ? [330,901.469, 1,614,779.25]
------------------------	----------------------------------	--

Feature column name	Column ID	Data type	Status ↓	Value
bedrooms	603580206602518528	Numeric	Required	4
condition	6584360511750537216	Categorical	Required	4
sqft_living	4278517502536843264	Numeric	Required	3000
waterfront	1972674493323149312	Categorical	Required	0

☐ Generate feature importance

PREDICT RESET

Predicted price = 771,622.313

95% prediction Interval: \$330,901.47 - \$1,614,779.25